

## Assignment-IV

### Fertilizer recommendation system for disease prediction

Date	26 October 2022
Student name	R Vaishnavi
Student roll no	621519205050
Maximum marks	2 marks

```
import pandas as pd import numpy as np import
matplotlib.pyplot as plt import seaborn as sns from
sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder from
tensorflow.keras.models import Model
from tensorflow.keras.layers import LSTM,Activation, Dense, Dropout, Input,
Embedding
from tensorflow.keras.optimizers import RMSprop from
tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing import sequence from
tensorflow.keras.utils import to_categorical from
tensorflow.keras.callbacks import EarlyStopping
%matplotlib inline import
csv
```

```
with open('/spam.csv', 'r') as csvfile: reader
= csv.reader(csvfile) df =
pd.read_csv(r'/spam.csv',encoding='latin-1')
df.head()
```

```
      v1                                     v2 Unnamed: 2  \ 0
ham  Go until jurong point, crazy.. Available only ...      NaN
1    ham                                Ok lar... Joking wif u oni...      NaN
2    spam  Free entry in 2 a wkly comp to win FA Cup fina...      NaN
3    ham  U dun say so early hor... U c already then say...      NaN    4
      ham  Nah I don't think he goes to usf, he lives aro...      NaN
```

```
Unnamed: 3 Unnamed: 4
0      NaN      NaN
1      NaN      NaN
2      NaN      NaN
```

```

3      NaN      NaN  4      NaN      NaN  df.drop(['Unnamed: 2',
      'Unnamed: 3', 'Unnamed: 4'],axis=1,inplace=True) df.info()

```

```


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    v1      5572 non-null    object
1    v2      5572 non-null    object
dtypes: object(2) memory usage:
87.2+ KB sns.countplot(df.v1)

```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43:

FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation. FutureWarning

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f5197dac250>



```

X = df.v2 Y = df.v1
le = LabelEncoder() Y =
le.fit_transform(Y)
Y = Y.reshape(-1,1)

```

```

X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.20)

```

```

max_words = 1000 max_len
= 150
tok = Tokenizer(num_words=max_words) tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train) sequences_matrix =
sequence.pad_sequences(sequences,maxlen=max_len)

def RNN():
    inputs = Input(name='inputs',shape=[max_len])
    layer = Embedding(max_words,50,input_length=max_len)(inputs)
    layer = LSTM(128)(layer) layer =
    Dense(256,name='FC1')(layer) layer =
    Activation('relu')(layer) layer = Dropout(0.5)(layer)
    layer = Dense(1,name='out_layer')(layer) layer =
    Activation('tanh')(layer) model =
    Model(inputs=inputs,outputs=layer) return model

model = RNN() model.summary()
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy',
'mse','mae'])

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
inputs (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 128)	91648
FC1 (Dense)	(None, 256)	33024
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0

```

=====
Total params: 174,929
Trainable params: 174,929
Non-trainable params: 0

```

```

model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,

validation_split=0.2,callbacks=[EarlyStopping(monitor='val_loss',min_delta=0.
0001)])

```

Epoch 1/10

```
28/28 [=====] - 17s 486ms/step - loss: 0.2960 -  
accuracy: 0.8819 - mse: 0.0821 - mae: 0.1563 - val_loss: 0.1341 -  
val_accuracy: 0.9675 - val_mse: 0.0344 - val_mae: 0.1237 Epoch 2/10  
28/28 [=====] - 13s 462ms/step - loss: 0.1149 -  
accuracy: 0.9764 - mse: 0.0381 - mae: 0.1538 - val_loss: 0.1321 -  
val_accuracy: 0.9798 - val_mse: 0.0437 - val_mae: 0.1695
```

```
<keras.callbacks.History at 0x7f5193192590>
```

```
test_sequences = tok.texts_to_sequences(X_test) test_sequences_matrix =  
sequence.pad_sequences(test_sequences,maxlen=max_len) accr =  
model.evaluate(test_sequences_matrix,Y_test)
```

```
35/35 [=====] - 3s 78ms/step - loss: 0.1590 -  
accuracy: 0.9812 - mse: 0.0451 - mae: 0.1733
```

```
print('Test set\n Loss: {:.3f}\n Accuracy:  
{:.3f}'.format(accr[0],accr[1]))
```

```
Test set
```

```
Loss: 0.159 Accuracy: 0.981
```

```
model.save("./assign4model.h5")
```

```
from tensorflow.keras.models import load_model m2  
= load_model("./assign4model.h5")  
m2.evaluate(test_sequences_matrix,Y_test)
```

```
35/35 [=====] - 3s 68ms/step - loss: 0.1590 -  
accuracy: 0.9812 - mse: 0.0451 - mae: 0.1733
```

```
[0.1589982509613037,  
0.9811659455299377,  
0.04506031796336174,  
0.17333826422691345]
```